A concurrent multi-threaded implementation of Conway's Game of Life

William Evans (qv18211), William Hannigan (yz18341)

Team 45, University of Bristol, United Kingdom

3 December 2019

Abstract

A comparison of a concurrent implementation of Conway's Game of Life over a single-threaded implementation. In this report, we aim to utilise goroutines to create a concurrent, multi-threaded Golang program simulating Conway's Game of Life on an image matrix to achieve a faster implementation over a single-threaded program. We use CPU profiling to analyse the bottlenecks within different implementations of this simulation. Experimental results show that our concurrent implementation produces a speed decrease of 313.5% over our single-threaded implementation.

Introduction

Concurrent, parallel computation can allow programs to be made more efficient by splitting and distributing work so that it can be processed simultaneously. Traditional multi-threaded programs rely on the use of system threads to allow for concurrent, parallel computation. These traditional threads have problems due to the fact that data is communicate between threads through the use of shared memory. Utilising shared memory requires the use of features such as semaphores and mutual-exclusion locks (Lee, 2006) which lead to a slow down of processing. This implementation has been written in Golang. Golang is a concurrent, garbage-collected programming language (Golang.org, n.d.).

By writing this implementation in Golang, we are able to use Communicating sequential processes (*CSP*) style programming aiming to achieve more efficient concurrent, parallel computation. Through the use of goroutines and channels instead of traditional threads Golang.org. (n.d.) we utilise CSP style programming to reduce the amount of time

our program spends locking and unlocking shared memory. Goroutines operate by being moved around between a set amount of system threads by the scheduler when one thread is blocked by a blocking operation in a goroutine. This method allows for concurrent applications to be built in a functional manner. By creating a concurrent, parallel implementation we aim to provide a performance increase over our single-threaded implementation by reducing the time it takes to process each turn of the Game of Life on an image.

Functionality and Design

Our implementation first reads in the initial Game of Life board from a PGM image and sends this to a main distributor goroutine. The distributor goroutine is responsible for splitting up the image into n parts and communicating it to n workers byte-by-byte over a buffered channel for each worker, with worker each running as a separate goroutine. These separate workers operate on their part of the image and then communicate the new version of the image back to the distributor (also byte-by-byte over a buffered channel for each worker) where it is then reconstructed at the end of each turn in the distributor.

The image (w x h), where w is the width h is the height, is split into n parts. Each part has equal width $pW = w$ and are of equal height, if n is a power of 2 where the height of a part is $pH = \lfloor \frac{h}{n} \rfloor$. If n is not a power of 2, all parts are of equal height pH except part n which has height: $pH = h - ((n-1) \times \lfloor \frac{h}{n} \rfloor)$, allowing the image to processed on any n workers where n is a multiple of 2.

Alongside this, user control has been implemented through the use of another *control* goroutine. This goroutine receives keypresses

using the termbox function PollEvent() (Reinke, 2019). Keypresses are then sent to a key channel. At the end of each turn, the distributor attempts to read from the key channel and if a keypress is present, reacts in the following manner. Upon the letter q being pressed, processing is stopped and a final image is outputted depicting the final turn that has been fully processed. This is done using the *writePgmImage* function provided in pgm.go in the skeleton code. The world is sent byte-by-byte to the PGM goroutine using a channel. Upon the letter p, processing is paused until p is pressed again. Upon the letter s, an image is outputted depicting the current turn in the same way as the letter q but the program is not stopped.

A *ticker* goroutine has also been implemented which receives commands on different channels. The distributor sends the number of alive cells at the end of each turn on a *numAliveCells* channel which is received by the ticker goroutine. The number of alive cells is printed every 2 seconds through the use of a ticker provided by Golang's time package (Golang.org, n.d.). Upon receiving a *p* keypress, the printing is paused by sending *true* on the ticker's pause channel. This is then resumed by sending *false* on the ticker's pause channel when another *p* keypress is received.

Through the use of channels and sending the image byte-by-byte on each channel, memory sharing is never used and therefore this concurrent implementation is thread safe. The current stage of this implementation solves the problem of executing our program on multiple cores in comparison to the initial single-threaded implementation which only makes use of a single system core.

Experiments

We demonstrate the performance of both our single-threaded implementation and our concurrent, multi-threaded implementation by executing both implementations on the University's lab machines with the CPU profiler and benchmarking tools provided by Go-lang's tools (Golang.org, n.d.). By using the Golang benchmarking tool and the benchmarks provided in the skeleton code, we were able to measure the time to execute 1000 turns for each implementation on a varying number of threads. This was all tested on the 128x128 image provided in the skeleton code.

Table 1: Runtimes for 1000 turns of 128x128 when using 2, 4 and 8 threads

Implementation	Runtime (seconds)		
	n=2	n=4	n=8
Single threaded	0.720	0.719	0.718
Baseline	0.735	0.528	0.397
Stage 3 multithreaded	2.98	4.67	8.28

Table 1 shows the results of our runtime experiment. Each implementation was executed three times on the 128x128 image. The average runtime of all these executions was taken as the result to ensure that a fair test was conducted.

Table 2: CPU usage for 1000 turns of 128x128 when using 2, 4 and 8 threads

Implementation	CPU usage (percentage)		
	n=2	n=4	n=8
Single threaded	104	104	104
Baseline	185	296	394
Stage 3 multithreaded	145	172	203

Table 2 shows the results of our CPU usage experiment. Each implementation was executed three times on the 128x128 image and the CPU usage was calculated using Linux's built in time tool returning a percentage, these commands were provided in the skeleton code. The average CPU usage of all these executions was taken as the result to ensure that a fair test was conducted.

Using the CPU profiler we analysed which top 4 functions requiring the largest

2

amount of CPU time within the single-threaded and Stage 3 multithreaded implementation of our program. This was conducted using the *make cpuprofile* command provided in the skeleton code which outputs a CPU profile. This was then analysed using Golang's internal tool *go tool pprof cpu.prof* followed by the command *top*.

Table 3: Top 4 functions ordered by most CPU usage to least

Implementation	Functions			
	1st	**2nd**	**3rd**	**4th**
Single threaded	getNum LiveNei ghbours	getNei ghbour LifeVal ue	distrib utor	getNe wLife Value
Stage 3 multithreaded	lock	unlock	chanre cv	chans end

Table 4: Time to process input images using our Stage 3 implementation

Image	Runtime (seconds)
16x16	0.0661
64x64	0.764
128x128	3.06
256x256	11.3
512x512	44.2

All tests were conducted on the same lab machine, when only one user was logged in to ensure that differences shown are just due to differences in implementation and not environmental differences.

Results

Using Table 1, we can see that for n=2 threads, our single-threaded implementation outperforms all other implementations including the baseline (2.04%). We expect this to be due to how we calculate the number of alive neighbours being slightly more efficient than the baseline. As the number of threads increases we see that the single-threaded im-plementation's runtime stays almost constant. This is expected as it only runs on one CPU core and is not parallelised, shown by the CPU usage staying constant at 104% in Table 2.

As the number of threads increases to n=4, using Table 1, we see that the baseline implementation now outperforms all other implementations in runtime, with a 28.2% decrease in runtime. This is expected in comparison to the single-threaded implementation. As from Table 2 we can see that the baseline implementation now uses multiple cores, with a 60.0% increase in CPU utilisation. However our Stage 3 multithreaded implementation has a 56.7% increase in runtime despite utilising 18.6% more CPU.

From Table 1, we see that our multithreaded version has a runtime that is 1990% higher than the baseline implementation and 1050% higher than our single threaded implementation.

From Table 3, we see that this is most likely due to the fact that a lot of time is spent in internal runtime functions, specifically *runtime.lock, runtime.unlock, runtime.chanrecv and runtime.chansend.* All these functions are internal functions utilised by channels. This is due to the fact that the whole world is communicated with each worker byte-by-byte upon each turn. This means a lot of time is spent just reconstructing the world during each term in both the workers and distributors. This, in comparison to the amount of time that is spent calculating the life value of each cell in the world in the single-threaded version, means that our multithreaded version is much less efficient than our single threaded version despite being made concurrent.

From Table 4, we can see that for our Stage 3 multithreaded implementation, the runtime increases massively as the dimensions of the image increase.

Conclusion

To conclude, our results show that at this stage of our implementation, it is more effi-

cient to use our single-threaded implementation than our multi-threaded implementation due to the large amount of data we are transferring over channels.

To improve upon this and get a more efficient, parallelised implementation we could use a halo exchange scheme or just communicate the currently alive cells to and from the workers. This would reduce the amount of data transferred and thus reduce the amount of locking and unlocking done by channels.

An alternative method could also be to share memory with the workers by sharing the world slice, however this would not allow for the implementation to be used on multiple networked machines in the future.

Our multi-threaded implementation does however improve on our single-threaded one by introducing the use of ticker and control threads which give the user more control when running the program and analysing what is going on within the Game of Life. It also sets up some necessary features, such as input and output channels for workers which would be used in later stages of the project.

References

[1] Lee, E. (2006). The Problem with Threads. Computer, 39(5), pp.33-42.
[2] Golang.org. (n.d.). The Go Programming Language Specification - The Go Programming Language. [online] Available at: https://golang.org/ref/spec#Introduction [Accessed 3 Dec. 2019].
[3] Golang.org. (n.d.). Why goroutines instead of threads? - Design FAQ - The Go Programming Language. [online] Available at: https://golang.org/doc/faq#csp [Accessed 3 Dec. 2019].
[4] Reinke, G. (2019). nsf/termbox-go. [online] GitHub. Available at: https://github.com/nsf/termbox-go [Accessed 3 Dec. 2019].
[5] Golang.org. (n.d.). time - The Go Programming Language. [online] Available at: https://golang.org/pkg/time/#NewTicker [Accessed 3 Dec. 2019].
[6] Golang.org. (n.d.). Diagnostics - The Go Programming Language. [online] Available at: https://golang.org/doc/diagnostics.html#profiling [Accessed 3 Dec. 2019].